

Instantaneous Time-Step Model

Dave Lorenz

December 3, 2015

This example illustrates how to set up and use a instantaneous time-step model. These models are typically used when there is additional explanatory variable information such as surrogate unit values, like specific conductance. The intent is often to model both the concentration or flux at any time and the load over a period of time.

This example uses data from the Bad River near Odanah, Wisc., USGS gaging station 04027000. The example will build a model of chloride.

```
> # Load the necessary packages and the data
> library(rloadest)
> library(dataRetrieval)
> # What unit values are available?
> subset(whatNWISdata("04027000", "uv"),
+   select=c("parm_cd", "srsname", "begin_date", "end_date"))
```

	parm_cd		srsname	begin_date	end_date
9	00010	Temperature, water		2011-03-03	2015-12-03
21	00060	Stream flow, mean. daily		2007-10-01	2015-12-03
26	00065	Height, gage		2015-08-05	2015-12-03
32	00095	Specific conductance		2011-03-06	2015-12-03
38	00300	Oxygen		2011-03-03	2015-12-03
44	00400	pH		2011-03-17	2015-12-03
282	63680	Turbidity		2011-03-17	2015-12-03

```
> # Get the QW data
> BadQW <- importNWISqw("04027000", "00940",
+   begin.date="2011-04-01", end.date="2014-09-30")
> # Merge data and time and set timezone (2 steps)
> BadQW <- transform(BadQW, dateTime=sample_dt + as.timeDay(sample_tm))
> BadQW <- transform(BadQW, dateTime=setTZ(dateTime, tzone_cd))
> # Now the Unit values data
> BadUV <- readNWISuv("04027000", c("00060", "00095", "00300", "63680"),
```

```

+   startDate="2011-04-01", endDate="2014-09-30", tz="America/Chicago")
> BadUV <- renameNWISColumns(BadUV)
> names(BadUV)

[1] "agency_cd"      "site_no"      "dateTime"      "tz_cd"
[5] "Flow_Inst"      "SpecCond_Inst" "DO_Inst"       "Turb_Inst"
[9] "Flow_Inst_cd"   "SpecCond_Inst_cd" "DO_Inst_cd"    "Turb_Inst_cd"

> # Strip _Inst off column names
> names(BadUV) <- sub("_Inst", "", names(BadUV))
> # Merge the data
> BadData <- mergeNearest(BadQW, "dateTime", right=BadUV, dates.right="dateTime",
+   max.diff="4 hours")
> # Rename the left-hand dateTime column
> names(BadData)[9] <- "dateTime"

```

1 Build the Instantaneous Time-Step Model

The first step in building the model is to determine which of the surrogates are most appropriate to include in the model. There can be many factors that contribute to deciding which explanatory variables to include in the model. From previous experience the user may decide to include or exclude specific surrogates and flow or seasonal terms. For this example, temperature (parameter code 00010) and pH (parameter code 000400) were excluded as they typically have very little influence on nitrate or nitrate concentration. Other factors include the availability of surrogate values. The output in the code below indicates that NTU_Turb has few observations (more missing values) that Turb, and will not be included in the candidate explanatory variables.

```
> # Print the number of missing values in each column
> sapply(BadData, function(col) sum(is.na(col)))
```

site_no.left	sample_dt	sample_tm	tzzone_cd	medium_cd	sample_end_dt
0	0	0	0	0	0
sample_end_tm	Chloride	dateTime	agency_cd	site_no.right	dateTime.right
0	0	0	0	0	0
tz_cd	Flow	SpecCond	DO	Turb	Flow_cd
0	22	21	15	26	22
SpecCond_cd	DO_cd	Turb_cd			
21	15	26			

This example will include the other surrogates and flow and seasonal terms in the candidate model. The code below demonstrates the use of `selBestSubset` to select the initial candidate model.

```
> # Create the and print the candidate model.
> BadChloride.lm <- selBestSubset(Chloride ~ log(Flow) + fourier(dateTime) +
+   log(SpecCond) + log(DO) + log(Turb), data=BadData,
+   flow="Flow", dates="dateTime", time.step="instantaneous",
+   station="Bad River near Odanah", criterion="SPCC")
> print(BadChloride.lm)
```

*** Load Estimation ***

Station: Bad River near Odanah
Constituent: Chloride

Number of Observations: 85
Number of Uncensored Observations: 85

Center of Decimal Time: 2012.58

Center of ln(Q): 6.772

Period of record: 2011-04-12 07:59:00 to 2014-07-15 14:15:00

Model Evaluation Criteria Based on AMLE Results

```
-----  
          Step Df Deviance Resid. Df Resid. Dev  SPCC  
1          NA      NA      77      -15.95 19.59  
2 - log(D0)  1    0.3468      78      -15.60 15.50  
Model # 99 selected
```

Selected Load Model:

Chloride ~ log(Flow) + fourier(dateTime) + log(SpecCond) + log(Turb)

Model coefficients:

	Estimate	Std. Error	z-score	p-value
(Intercept)	-5.5744	1.45776	-3.824	0.0001
log(Flow)	1.2973	0.09535	13.606	0.0000
fourier(dateTime)sin(k=1)	0.3669	0.09069	4.045	0.0001
fourier(dateTime)cos(k=1)	0.2045	0.07107	2.877	0.0036
log(SpecCond)	1.3980	0.22357	6.253	0.0000
log(Turb)	-0.2498	0.05468	-4.569	0.0000

AMLE Regression Statistics

Residual variance: 0.05243

R-squared: 95.56 percent

G-squared: 264.7 on 5 degrees of freedom

P-value: <0.0001

Prob. Plot Corr. Coeff. (PPCC):

r = 0.9659

p-value = 5e-04

Serial Correlation of Residuals: 0.2815

Variance Inflation Factors:

	VIF
log(Flow)	23.010
fourier(dateTime)sin(k=1)	3.568
fourier(dateTime)cos(k=1)	1.203
log(SpecCond)	11.621
log(Turb)	10.351

Comparison of Observed and Estimated Loads

```

Summary Stats: Loads in kg/d
-----
      Min  25%   50%   75%   90%   95%   Max
Est 920 7700 11500 21400 28400 35000 58000
Obs 920 7260 11500 21000 32800 35200 40400

```

Bias Diagnostics

```

-----
Bp: 0.08092 percent
PLR: 1.001
E: 0.8249

```

Only log(DO) was dropped from the model. The printed report indicates some potential problems with the regression—the PPCC test indicates the residuals are not normally distributed and several variance inflation factors are relatively large, greater than 10. But the bias diagnostics show very little bias in the comparison of the estimated to observed values.

A few selected graphs will help understand the issues identified in the printed report and suggest an alternative model. Figure 1 shows the residuals versus fitted graph, which indicates some very large residuals at larger fitted values. It also suggests some heteroscedasticity in the residual pattern.

```

> # Plot the overall fit, choose plot number 2.
> setSweave("graph01", 6, 6)
> plot(BadChloride.lm, which = 2, set.up=FALSE)
> dev.off()

```

```

null device
      1

```




Figure 1. The residuals versus fitted graph.

The S-L plot is not shown. The residual Q-normal graph indicates the reason for the very low p-value indicated by the PPCC test—the large residual values indicated in figure 1 skew the distribution.

```
> # Plot the residual Q-normal graph.
> setSweave("graph02", 6, 6)
> plot(BadChloride.lm, which = 5, set.up=FALSE)
> dev.off()
```

null device

1




Figure 2. The residual Q-normal graph.

A complete review of the partial residual graphs is not included in this example. Only the partial residual for $\log(\text{Turb})$ is shown. The graph indicates the lack of fit, especially for the largest values of Turbidity. This suggests that the log transform is not appropriate.

```
> # Plot the residual Q-normal graph.
> setSweave("graph03", 6, 6)
> plot(BadChloride.lm, which = "log(Turb)", set.up=FALSE)
> dev.off()
```

```
null device
1
```




Figure 3. The partial residual for $\log(\text{Turb})$ graph.

Build the model excluding $\log(\text{DO})$ that was dropped in the subset selection procedure and changing $\log(\text{Turb})$ to Turb .

```
> # Create the and print the revised model.
> BadChloride.lr <- loadReg(Chloride ~ log(Flow) + fourier(dateTime) +
+   log(SpecCond) + Turb, data=BadData,
+   flow="Flow", dates="dateTime", time.step="instantaneous",
+   station="Bad River near Odanah")
> print(BadChloride.lr, load.only=FALSE)
```

*** Load Estimation ***

Station: Bad River near Odanah
Constituent: Chloride

Number of Observations: 85
Number of Uncensored Observations: 85
Center of Decimal Time: 2012.58
Center of $\ln(Q)$: 6.772

Period of record: 2011-04-12 07:59:00 to 2014-07-15 14:15:00

Selected Load Model:

Chloride ~ log(Flow) + fourier(dateTime) + log(SpecCond) + Turb

Model coefficients:

	Estimate	Std. Error	z-score	p-value
(Intercept)	-7.33549	1.0913995	-6.721	0e+00
log(Flow)	1.27900	0.0550056	23.252	0e+00
fourier(dateTime)sin(k=1)	0.24883	0.0681555	3.651	3e-04
fourier(dateTime)cos(k=1)	0.29203	0.0570742	5.117	0e+00
log(SpecCond)	1.65230	0.1687211	9.793	0e+00
Turb	-0.00122	0.0001442	-8.464	0e+00

AMLE Regression Statistics

Residual variance: 0.03476

R-squared: 97.05 percent

G-squared: 299.6 on 5 degrees of freedom

P-value: <0.0001

Prob. Plot Corr. Coeff. (PPCC):

r = 0.9916

p-value = 0.2894

Serial Correlation of Residuals: 0.166

Variance Inflation Factors:

	VIF
log(Flow)	11.549
fourier(dateTime)sin(k=1)	3.039
fourier(dateTime)cos(k=1)	1.170
log(SpecCond)	9.982
Turb	3.371

Comparison of Observed and Estimated Loads

Summary Stats: Loads in kg/d

	Min	25%	50%	75%	90%	95%	Max
Est	917	7530	11900	19300	29400	36400	60900
Obs	920	7260	11500	21000	32800	35200	40400

Bias Diagnostics

Bp: 0.4108 percent

PLR: 1.004

E: 0.8596

Selected Concentration Model:

Chloride ~ log(Flow) + fourier(dateTime) + log(SpecCond) + Turb

Model coefficients:

	Estimate	Std. Error	z-score	p-value
(Intercept)	-8.23018	1.0913995	-7.541	0e+00
log(Flow)	0.27900	0.0550056	5.072	0e+00
fourier(dateTime)sin(k=1)	0.24883	0.0681555	3.651	3e-04
fourier(dateTime)cos(k=1)	0.29203	0.0570742	5.117	0e+00
log(SpecCond)	1.65230	0.1687211	9.793	0e+00
Turb	-0.00122	0.0001442	-8.464	0e+00

AMLE Regression Statistics

Residual variance: 0.03476

R-squared: 76.28 percent

G-squared: 122.3 on 5 degrees of freedom

P-value: <0.0001

Prob. Plot Corr. Coeff. (PPCC):

r = 0.9916

p-value = 0.2894

Serial Correlation of Residuals: 0.166

Comparison of Observed and Estimated Concentrations

Summary Stats: Concentrations in mg/l

	Min	25%	50%	75%	90%	95%	Max
Est	1.15	2.00	2.46	2.91	3.88	4.10	5.25
Obs	0.76	1.85	2.37	3.08	3.96	4.28	4.96

Bias Diagnostics

Bp: 0.3245 percent

PCR: 1.003

E: 0.8115

The report for the revised model indicates less severe problems than from the first candidate model—the p-value for the PPCC test is greater than 0.05, the variance inflation factors are lower although `log(Flow)` is still greater than 10, and the bias diagnostics from the observed and estimated loads and concentrations are still good.

A review of selected diagnostic plots indicates a much better overall fit. Figure 4 shows the residuals versus fitted graph, which indicates a less severe problem of large residuals at larger fitted values. It also suggests some heteroscedasticity in the residual pattern as with the first candidate model.

```
> # Plot the overall fit, choose plot number 2.
> setSweave("graph04", 6, 6)
> plot(BadChloride.lm, which = 2, set.up=FALSE)
> dev.off()
```

```
null device
      1
```




Figure 4. The residuals versus fitted graph for the revised model.

For this model, the S-L plot is shown. It shows an increase in heteroscedasticity as the fitted values increase. That heteroscedasticity can introduce bias into the estimated values as the bias correction factor will be a bit too small for the larger values and too large for the smaller values. The potential bias for this model is expected to be small because the residual variance is small, 0.03476 natural log units, therefore the bias correction is

very small, less than 2 percent, and the potential change to the bias correction very small, much less than 1/2 percent.

```
> # Plot the S-L graph.
> setSweave("graph05", 6, 6)
> plot(BadChloride.lmr, which = 3, set.up=FALSE)
> dev.off()
```

```
null device
      1
```




Figure 5. The S-L graph for the revised model.

The residual Q-normal graph shows much better agreement to the normal distribution than the original candidate model—the effect of the lowest residuals is much less.

```
> # Plot the residual Q-normal graph.
> setSweave("graph06", 6, 6)
> plot(BadChloride.lmr, which = 5, set.up=FALSE)
> dev.off()
```

```
null device
      1
```




Figure 6. The residual Q-normal graph for the revised model.

A complete review of the partial residual graphs is not included in this example. Only the partial residual for **Turb** is shown to compare to the original model. In this case, the untransformed variable appears to fit reasonably well.

```
> # Plot the residual Q-normal graph.
> setSweave("graph07", 6, 6)
> plot(BadChloride.lm, which = "Turb", set.up=FALSE)
> dev.off()
```

```
null device
      1
```




Figure 7. The partial residual for Turb graph for the revised model.

2 Instantaneous Concentrations

Estimating the instantaneous concentrations or loads from the model is relatively straight forward. The `predConc` and `predLoad` functions will estimate concentrations or loads, actually fluxes, for any time, given explanatory variables with no missing values. This example will focus on a single day, June 30, 2014.

```
> # Extract one day from the UV data
> Bad063014 <- subset(BadUV, as.Date(as.POSIXlt(dateTime)) == "2014-06-30")
> # Remove the unnecessary surrogates from the data set.
> # This reduces the likelihood of missing values in the dataset
> Bad063014 <- Bad063014[, c("dateTime", "Flow", "SpecCond", "Turb")]
> # Simple check
> any(is.na(Bad063014))
```

```
[1] FALSE
```

```
> # Estimate concentrations
> Bad063014.est <- predConc(BadChloride.lmr, Bad063014, by="unit")
> # Display the first and last few rows.
> head(Bad063014.est)
```

	Date	Flow	Conc	Std.Err	SEP	L95	U95
1	2014-06-30 00:00:00	909	3.315789	0.1483089	0.6410364	2.223455	4.766606
2	2014-06-30 00:15:00	909	3.299546	0.1462430	0.6375876	2.213003	4.742480
3	2014-06-30 00:30:00	909	3.311449	0.1477848	0.6401211	2.220654	4.760175
4	2014-06-30 00:45:00	909	3.311294	0.1477819	0.6400920	2.220548	4.759954
5	2014-06-30 01:00:00	903	3.305057	0.1468059	0.6387257	2.216595	4.750585
6	2014-06-30 01:15:00	903	3.249339	0.1411140	0.6272261	2.180268	4.668656

```
> tail(Bad063014.est)
```

	Date	Flow	Conc	Std.Err	SEP	L95	U95
91	2014-06-30 22:30:00	770	3.058103	0.1205894	0.5876831	2.055698	4.387281
92	2014-06-30 22:45:00	770	3.046816	0.1194058	0.5853628	2.048326	4.370707
93	2014-06-30 23:00:00	770	3.057821	0.1205899	0.5876313	2.055505	4.386882
94	2014-06-30 23:15:00	764	3.043614	0.1190648	0.5847037	2.046236	4.366004
95	2014-06-30 23:30:00	764	3.050892	0.1198473	0.5862035	2.050985	4.376700
96	2014-06-30 23:45:00	758	2.988006	0.1148736	0.5736139	2.009432	4.285213

```
> # The daily mean concentration can also be easily estimated
> predConc(BadChloride.lmr, Bad063014, by="day")
```

	Date	Flow	Conc	Std.Err	SEP	L95	U95
1	2014-06-30	829.4792	3.158582	0.1302371	0.6083021	2.121373	4.534724

```
> # Compare to the mean of the unit values:  
> with(Bad063014.est, mean(Conc))
```

```
[1] 3.158582
```


3 Aggregate Loads

Estimating concentrations or loads by day assumes, but does not require consistent number of unit values per day. Both `predLoad` and `predConc` assume that inconsistent number of unit values per day are due to missing values and return missing values for the estimates for days that do not have the average number of observations per day. Inconsistent number of observations per day can be the result of deleted bad values, maintenance, or a change in frequency of sampling. The data can be resampled to a uniform number per day using the `resampleUVdata` function or the check can be suppressed by setting the `allow.incomplete` argument to `TRUE`.

Estimating loads for periods longer than one day requires consistent number of unit values in each day. The consistent number per day is required to be able to keep track of within-day and between day variances. The `resampleUVdata` function can be used to force a consistent number of unit values per day. It is not required for this example, but useful when the unit values are not consistent or when there is a change to or from daylight savings time.

Just as with estimating instantaneous values, missing values are not permitted. Missing values can occur with surrogates due to short-term malfunctions, calibration, or long-term malfunctions. Missing values from short-term malfunctions, generally spikes in the data that are removed during processing, or that occur during calibrations can easily be interpolated using the `fillMissing` function in (`smwrBase`) and are illustrated in this example. Longer-term missing values are much more difficult to fix. They require the careful balancing of need, developing alternate regression models and possible caveats of the interpretation of loads.

```
> # Extract one month from the UV data, done in two steps
> Bad0714 <- subset(BadUV, as.Date(as.POSIXlt(dateTime)) >= "2014-07-01")
> Bad0714 <- subset(Bad0714, as.Date(as.POSIXlt(dateTime)) <= "2014-07-31")
> # Remove the unnecessary surrogates from the data set.
> # This reduces the likelihood of missing values in the dataset
> Bad0714 <- Bad0714[, c("dateTime", "Flow", "SpecCond", "Turb")]
> # Simple check on each column, how many in each column?
> sapply(Bad0714, function(x) sum(is.na(x)))
```

dateTime	Flow	SpecCond	Turb
0	0	3	17

```
> # Fix each column, using the defaults of fillMissing
> Bad0714$SpecCond <- fillMissing(Bad0714$SpecCond)
> Bad0714$Turb <- fillMissing(Bad0714$Turb)
> # Verify filled values
> sapply(Bad0714, function(x) sum(is.na(x)))
```

```

dateTime      Flow SpecCond      Turb
      0          0          0          0

```

```

> # Estimate daily loads
> Bad0714.day <- predLoad(BadChloride.lr, Bad0714, by="day")
> # Display the first and last few rows.
> head(Bad0714.day)

```

	Date	Flow	Flux	Std.Err	SEP	L95	U95
1	2014-07-01	690.9271	4999.268	188.4703	961.0055	3360.170	7172.870
2	2014-07-02	733.5521	5389.557	209.8556	1041.0709	3615.312	7745.524
3	2014-07-03	702.1146	5062.956	194.9826	975.9416	3399.136	7271.018
4	2014-07-04	519.5417	3620.096	135.9578	696.7035	2432.021	5196.104
5	2014-07-05	399.8750	2895.400	112.0630	556.6967	1945.926	4154.565
6	2014-07-06	336.1875	2518.843	101.0843	484.4965	1692.567	3614.753

```

> tail(Bad0714.day)

```

	Date	Flow	Flux	Std.Err	SEP	L95	U95
26	2014-07-26	204.7604	1647.717	72.96417	318.4157	1105.095	2368.332
27	2014-07-27	200.3438	1723.719	76.54007	334.0974	1154.653	2480.076
28	2014-07-28	275.9167	2509.353	114.02758	486.0095	1681.434	3609.530
29	2014-07-29	277.8333	2379.386	105.39399	460.0444	1595.477	3420.586
30	2014-07-30	372.5625	3184.768	147.53720	619.2851	2130.508	4587.265
31	2014-07-31	477.9062	3849.708	183.75181	748.8520	2574.951	5545.704

```

> # And the month
> Bad0714.mon <- predLoad(BadChloride.lr, Bad0714, by="month")
> Bad0714.mon

```

	Period	Ndays	Flux	Std.Err	SEP	L95	U95
1	July 2014	31	3060.67	118.2118	115.4483	2840.608	3293.095

```

> # Compare to the results using the approximate standard error:
> # For long periods, the processing time to the exact seopt can be very large
> # and may be desireable to use the approximation.
> predLoad(BadChloride.lr, Bad0714, by="month", seopt="app")

```

	Period	Ndays	Flux	Std.Err	SEP	L95	U95
1	July 2014	31	3060.67	113.787	115.4478	2840.609	3293.094

```

> # Compare to the mean of the daily values:
> with(Bad0714.day, mean(Flux))

```

```

[1] 3060.67

```